Thrift: Scalable Cross-Language Services Implementation

Mark Slee, Aditya Agarwal and Marc Kwiatkowski

Facebook, 156 University Ave, Palo Alto, CA

{mcslee,aditya,marc}@facebook.com

Abstract

Thrift is a software library and set of code-generation tools developed at Facebook to expedite development and implementation of efficient and scalable backend services. Its primary goal is to enable efficient and reliable communication across programming languages by abstracting the portions of each language that tend to require the most customization into a common library that is implemented in each language. Specifically, Thrift allows developers to define datatypes and service interfaces in a single language-neutral file and generate all the necessary code to build RPC clients and servers.

This paper details the motivations and design choices we made in Thrift, as well as some of the more interesting implementation details. It is not intended to be taken as research, but rather it is an exposition on what we did and why.

1. Introduction

As Facebook's traffic and network structure have scaled, the resource demands of many operations on the site (i.e. search, ad selection and delivery, event logging) have presented technical requirements drastically outside the scope of the LAMP framework. In our implementation of these services, various programming languages have been selected to optimize for the right combination of performance, ease and speed of development, availability of existing libraries, etc. By and large, Facebook's engineering culture has tended towards choosing the best tools and implementations available over standardizing on any one programming language and begrudgingly accepting its inherent limitations.

Given this design choice, we were presented with the challenge of building a transparent, high-performance bridge across many programming languages. We found that most available solutions were either too limited, did not offer sufficient datatype freedom, or suffered from subpar performance. [1]

The solution that we have implemented combines a language-neutral software stack implemented across numerous programming languages and an associated code generation engine that transforms a simple interface and data definition language into client and server remote procedure call libraries. Choosing static code generation over a dynamic system allows us to create validated code that can be run without the need for any advanced introspective run-time type checking. It is also designed to be as simple as possible for the developer, who can typically define all the necessary data structures and interfaces for a complex service in a single short file.

Surprised that a robust open solution to these relatively common problems did not yet exist, we committed early on to making the Thrift implementation open source.

---

[1] See Appendix A for a discussion of alternative systems.

In evaluating the challenges of cross-language interaction in a networked environment, some key components were identified:

Types. A common type system must exist across programming languages without requiring that the application developer use custom Thrift datatypes or write their own serialization code. That is, a C++ programmer should be able to transparently exchange a strongly typed STL map for a dynamic Python dictionary. Neither programmer should be forced to write any code below the application layer to achieve this. Section 2 details the Thrift type system.

Transport. Each language must have a common interface to bidirectional raw data transport. The specifics of how a given transport is implemented should not matter to the service developer. The same application code should be able to run against TCP stream sockets, raw data in memory, or files on disk. Section 3 details the Thrift Transport layer.

Protocol. Datatypes must have some way of using the Transport layer to encode and decode themselves. Again, the application developer need not be concerned by this layer. Whether the service uses an XML or binary protocol is immaterial to the application code. All that matters is that the data can be read and written in a consistent, deterministic matter. Section 4 details the Thrift Protocol layer.

Versioning. For robust services, the involved datatypes must provide a mechanism for versioning themselves. Specifically, it should be possible to add or remove fields in an object or alter the argument list of a function without any interruption in service (or, worse yet, nasty segmentation faults). Section 5 details Thrift's versioning system.

Processors. Finally, we generate code capable of processing data streams to accomplish remote procedure calls. Section 6 details the generated code and TProcessor paradigm.

Section 7 discusses implementation details, and Section 8 describes our conclusions.

2. Types

The goal of the Thrift type system is to enable programmers to develop using completely natively defined types, no matter what programming language they use. By design, the Thrift type system does not introduce any special dynamic types or wrapper objects. It also does not require that the developer write any code for object serialization or transport. The Thrift IDL (Interface Definition Language) file is logically a way for developers to annotate their data structures with the minimal amount of extra information necessary to tell a code generator how to safely transport the objects across languages.

2.1 Base Types

The type system rests upon a few base types. In considering which types to support, we aimed for clarity and simplicity over abun-

dance, focusing on the key types available in all programming languages, ommitting any niche types available only in specific languages.

The base types supported by Thrift are:

- `bool` A boolean value, true or false
- `byte` A signed byte
- `i16` A 16-bit signed integer
- `i32` A 32-bit signed integer
- `i64` A 64-bit signed integer
- `double` A 64-bit floating point number
- `string` An encoding-agnostic text or binary string

Of particular note is the absence of unsigned integer types. Because these types have no direct translation to native primitive types in many languages, the advantages they afford are lost. Further, there is no way to prevent the application developer in a language like Python from assigning a negative value to an integer variable, leading to unpredictable behavior. From a design standpoint, we observed that unsigned integers were very rarely, if ever, used for arithmetic purposes, but in practice were much more often used as keys or identifiers. In this case, the sign is irrelevant. Signed integers serve this same purpose and can be safely cast to their unsigned counterparts (most commonly in C++) when absolutely necessary.

2.2 Structs

A Thrift struct defines a common object to be used across languages. A struct is essentially equivalent to a class in object oriented programming languages. A struct has a set of strongly typed fields, each with a unique name identifier. The basic syntax for defining a Thrift struct looks very similar to a C struct definition. Fields may be annotated with an integer field identifier (unique to the scope of that struct) and optional default values. Field identifiers will be automatically assigned if omitted, though they are strongly encouraged for versioning reasons discussed later.

2.3 Containers

Thrift containers are strongly typed containers that map to the most commonly used containers in common programming languages. They are annotated using the C++ template (or Java Generics) style. There are three types available:

- `list<type>` An ordered list of elements. Translates directly into an STL `vector`, Java `ArrayList`, or native array in scripting languages. May contain duplicates.
- `set<type>` An unordered set of unique elements. Translates into an STL `set`, Java `HashSet`, `set` in Python, or native dictionary in PHP/Ruby.
- `map<type1,type2>` A map of strictly unique keys to values Translates into an STL `map`, Java `HashMap`, PHP associative array, or Python/Ruby dictionary.

While defaults are provided, the type mappings are not explicitly fixed. Custom code generator directives have been added to substitute custom types in destination languages (i.e. `hash_map` or Google's sparse hash map can be used in C++). The only requirement is that the custom types support all the necessary iteration primitives. Container elements may be of any valid Thrift type, including other containers or structs.

```
struct Example {
  1:i32 number=10,
  2:i64 bigNumber,
```

```
  3:double decimals,
  4:string name="thrifty"
}
```

In the target language, each definition generates a type with two methods, `read` and `write`, which perform serialization and transport of the objects using a Thrift TProtocol object.

2.4 Exceptions

Exceptions are syntactically and functionally equivalent to structs except that they are declared using the `exception` keyword instead of the `struct` keyword.

The generated objects inherit from an exception base class as appropriate in each target programming language, in order to seamlessly integrate with native exception handling in any given language. Again, the design emphasis is on making the code familiar to the application developer.

2.5 Services

Services are defined using Thrift types. Definition of a service is semantically equivalent to defining an interface (or a pure virtual abstract class) in object oriented programming. The Thrift compiler generates fully functional client and server stubs that implement the interface. Services are defined as follows:

```
service <name> {
  <returntype> <name>(<arguments>)
    [throws (<exceptions>)]
  ...
}
```

An example:

```
service StringCache {
  void set(1:i32 key, 2:string value),
  string get(1:i32 key) throws (1:KeyNotFound knf),
  void delete(1:i32 key)
}
```

Note that `void` is a valid type for a function return, in addition to all other defined Thrift types. Additionally, an `async` modifier keyword may be added to a `void` function, which will generate code that does not wait for a response from the server. Note that a pure `void` function will return a response to the client which guarantees that the operation has completed on the server side. With `async` method calls the client will only be guaranteed that the request succeeded at the transport layer. (In many transport scenarios this is inherently unreliable due to the Byzantine Generals' Problem. Therefore, application developers should take care only to use the async optimization in cases where dropped method calls are acceptable or the transport is known to be reliable.)

Also of note is the fact that argument lists and exception lists for functions are implemented as Thrift structs. All three constructs are identical in both notation and behavior.

3. Transport

The transport layer is used by the generated code to facilitate data transfer.

3.1 Interface

A key design choice in the implementation of Thrift was to decouple the transport layer from the code generation layer. Though Thrift is typically used on top of the TCP/IP stack with streaming

sockets as the base layer of communication, there was no compelling reason to build that constraint into the system. The performance tradeoff incurred by an abstracted I/O layer (roughly one virtual method lookup / function call per operation) was immaterial compared to the cost of actual I/O operations (typically invoking system calls).

Fundamentally, generated Thrift code only needs to know how to read and write data. The origin and destination of the data are irrelevant; it may be a socket, a segment of shared memory, or a file on the local disk. The Thrift transport interface supports the following methods:

- `open` Opens the tranpsort
- `close` Closes the tranport
- `isOpen` Indicates whether the transport is open
- `read` Reads from the transport
- `write` Writes to the transport
- `flush` Forces any pending writes

There are a few additional methods not documented here which are used to aid in batching reads and optionally signaling the completion of a read or write operation from the generated code.

In addition to the above `TTransport` interface, there is a `TServerTransport` interface used to accept or create primitive transport objects. Its interface is as follows:

- `open` Opens the transport
- `listen` Begins listening for connections
- `accept` Returns a new client transport
- `close` Closes the transport

3.2 Implementation

The transport interface is designed for simple implementation in any programming language. New transport mechanisms can be easily defined as needed by application developers.

3.2.1 TSocket

The `TSocket` class is implemented across all target languages. It provides a common, simple interface to a TCP/IP stream socket.

3.2.2 TFileTransport

The `TFileTransport` is an abstraction of an on-disk file to a data stream. It can be used to write out a set of incoming Thrift requests to a file on disk. The on-disk data can then be replayed from the log, either for post-processing or for reproduction and/or simulation of past events.

3.2.3 Utilities

The Transport interface is designed to support easy extension using common OOP techniques, such as composition. Some simple utilites include the `TBufferedTransport`, which buffers the writes and reads on an underlying transport, the `TFramedTransport`, which transmits data with frame size headers for chunking optimization or nonblocking operation, and the `TMemoryBuffer`, which allows reading and writing directly from the heap or stack memory owned by the process.

4. Protocol

A second major abstraction in Thrift is the separation of data structure from transport representation. Thrift enforces a certain

messaging structure when transporting data, but it is agnostic to the protocol encoding in use. That is, it does not matter whether data is encoded as XML, human-readable ASCII, or a dense binary format as long as the data supports a fixed set of operations that allow it to be deterministically read and written by generated code.

4.1 Interface

The Thrift Protocol interface is very straightforward. It fundamentally supports two things: 1) bidirectional sequenced messaging, and 2) encoding of base types, containers, and structs.

```
writeMessageBegin(name, type, seq)
writeMessageEnd()
writeStructBegin(name)
writeStructEnd()
writeFieldBegin(name, type, id)
writeFieldEnd()
writeFieldStop()
writeMapBegin(ktype, vtype, size)
writeMapEnd()
writeListBegin(etype, size)
writeListEnd()
writeSetBegin(etype, size)
writeSetEnd()
writeBool(bool)
writeByte(byte)
writeI16(i16)
writeI32(i32)
writeI64(i64)
writeDouble(double)
writeString(string)
```

```
name, type, seq = readMessageBegin()
                  readMessageEnd()
name =            readStructBegin()
                  readStructEnd()
name, type, id =  readFieldBegin()
                  readFieldEnd()
k, v, size =      readMapBegin()
                  readMapEnd()
etype, size =     readListBegin()
                  readListEnd()
etype, size =     readSetBegin()
                  readSetEnd()
bool =            readBool()
byte =            readByte()
i16 =             readI16()
i32 =             readI32()
i64 =             readI64()
double =          readDouble()
string =          readString()
```

Note that every `write` function has exactly one `read` counterpart, with the exception of `writeFieldStop()`. This is a special method that signals the end of a struct. The procedure for reading a struct is to `readFieldBegin()` until the stop field is encountered, and then to `readStructEnd()`. The generated code relies upon this call sequence to ensure that everything written by a protocol encoder can be read by a matching protocol decoder. Further note that this set of functions is by design more robust than necessary. For example, `writeStructEnd()` is not strictly necessary, as the end of a struct may be implied by the stop field. This method is a convenience for verbose protocols in which it is cleaner to separate these calls (e.g. a closing `</struct>` tag in XML).

4.2 Structure

Thrift structures are designed to support encoding into a streaming protocol. The implementation should never need to frame or compute the entire data length of a structure prior to encoding it. This is critical to performance in many scenarios. Consider a long list of relatively large strings. If the protocol interface required reading or writing a list to be an atomic operation, then the implementation would need to perform a linear pass over the entire list before encoding any data. However, if the list can be written as iteration is performed, the corresponding read may begin in parallel, theoretically offering an end-to-end speedup of $(kN - C)$, where N is the size of the list, k the cost factor associated with serializing a single element, and C is fixed offset for the delay between data being written and becoming available to read.

Similarly, structs do not encode their data lengths a priori. Instead, they are encoded as a sequence of fields, with each field having a type specifier and a unique field identifier. Note that the inclusion of type specifiers allows the protocol to be safely parsed and decoded without any generated code or access to the original IDL file. Structs are terminated by a field header with a special STOP type. Because all the basic types can be read deterministically, all structs (even those containing other structs) can be read deterministically. The Thrift protocol is self-delimiting without any framing and regardless of the encoding format.

In situations where streaming is unnecessary or framing is advantageous, it can be very simply added into the transport layer, using the TFramedTransport abstraction.

4.3 Implementation

Facebook has implemented and deployed a space-efficient binary protocol which is used by most backend services. Essentially, it writes all data in a flat binary format. Integer types are converted to network byte order, strings are prepended with their byte length, and all message and field headers are written using the primitive integer serialization constructs. String names for fields are omitted - when using generated code, field identifiers are sufficient.

We decided against some extreme storage optimizations (i.e. packing small integers into ASCII or using a 7-bit continuation format) for the sake of simplicity and clarity in the code. These alterations can easily be made if and when we encounter a performance-critical use case that demands them.

5. Versioning

Thrift is robust in the face of versioning and data definition changes. This is critical to enable staged rollouts of changes to deployed services. The system must be able to support reading of old data from log files, as well as requests from out-of-date clients to new servers, and vice versa.

5.1 Field Identifiers

Versioning in Thrift is implemented via field identifiers. The field header for every member of a struct in Thrift is encoded with a unique field identifier. The combination of this field identifier and its type specifier is used to uniquely identify the field. The Thrift definition language supports automatic assignment of field identifiers, but it is good programming practice to always explicitly specify field identifiers. Identifiers are specified as follows:

```
struct Example {
  1:i32 number=10,
  2:i64 bigNumber,
  3:double decimals,
  4:string name="thrifty"
```

```
}
```

To avoid conflicts between manually and automatically assigned identifiers, fields with identifiers omitted are assigned identifiers decrementing from -1, and the language only supports the manual assignment of positive identifiers.

When data is being deserialized, the generated code can use these identifiers to properly identify the field and determine whether it aligns with a field in its definition file. If a field identifier is not recognized, the generated code can use the type specifier to skip the unknown field without any error. Again, this is possible due to the fact that all datatypes are self delimiting.

Field identifiers can (and should) also be specified in function argument lists. In fact, argument lists are not only represented as structs on the backend, but actually share the same code in the compiler frontend. This allows for version-safe modification of method parameters

```
service StringCache {
  void set(1:i32 key, 2:string value),
  string get(1:i32 key) throws (1:KeyNotFound knf),
  void delete(1:i32 key)
}
```

The syntax for specifying field identifiers was chosen to echo their structure. Structs can be thought of as a dictionary where the identifiers are keys, and the values are strongly-typed named fields.

Field identifiers internally use the i16 Thrift type. Note, however, that the TProtocol abstraction may encode identifiers in any format.

5.2 Isset

When an unexpected field is encountered, it can be safely ignored and discarded. When an expected field is not found, there must be some way to signal to the developer that it was not present. This is implemented via an inner isset structure inside the defined objects. (Isset functionality is implicit with a null value in PHP, None in Python and nil in Ruby.) Essentially, the inner isset object of each Thrift struct contains a boolean value for each field which denotes whether or not that field is present in the struct. When a reader receives a struct, it should check for a field being set before operating directly on it.

```
class Example {
 public:
  Example() :
    number(10),
    bigNumber(0),
    decimals(0),
    name("thrifty") {}

  int32_t number;
  int64_t bigNumber;
  double decimals;
  std::string name;

  struct __isset {
    __isset() :
      number(false),
      bigNumber(false),
      decimals(false),
      name(false) {}
    bool number;
    bool bigNumber;
    bool decimals;
```

```
      bool name;
  } __isset;
...
}
```

5.3 Case Analysis

There are four cases in which version mismatches may occur.

1. *Added field, old client, new server.* In this case, the old client does not send the new field. The new server recognizes that the field is not set, and implements default behavior for out-of-date requests.

2. *Removed field, old client, new server.* In this case, the old client sends the removed field. The new server simply ignores it.

3. *Added field, new client, old server.* The new client sends a field that the old server does not recognize. The old server simply ignores it and processes as normal.

4. *Removed field, new client, old server.* This is the most dangerous case, as the old server is unlikely to have suitable default behavior implemented for the missing field. It is recommended that in this situation the new server be rolled out prior to the new clients.

5.4 Protocol/Transport Versioning

The `TProtocol` abstractions are also designed to give protocol implementations the freedom to version themselves in whatever manner they see fit. Specifically, any protocol implementation is free to send whatever it likes in the `writeMessageBegin()` call. It is entirely up to the implementor how to handle versioning at the protocol level. The key point is that protocol encoding changes are safely isolated from interface definition version changes.

Note that the exact same is true of the `TTransport` interface. For example, if we wished to add some new checksumming or error detection to the `TFileTransport`, we could simply add a version header into the data it writes to the file in such a way that it would still accept old log files without the given header.

6. RPC Implementation

6.1 TProcessor

The last core interface in the Thrift design is the `TProcessor`, perhaps the most simple of the constructs. The interface is as follows:

```
interface TProcessor {
  bool process(TProtocol in, TProtocol out)
    throws TException
}
```

The key design idea here is that the complex systems we build can fundamentally be broken down into agents or services that operate on inputs and outputs. In most cases, there is actually just one input and output (an RPC client) that needs handling.

6.2 Generated Code

When a service is defined, we generate a `TProcessor` instance capable of handling RPC requests to that service, using a few helpers. The fundamental structure (illustrated in pseudo-C++) is as follows:

```
Service.thrift
 => Service.cpp
    interface ServiceIf
    class ServiceClient : virtual ServiceIf
```

```
      TProtocol in
      TProtocol out
    class ServiceProcessor : TProcessor
      ServiceIf handler

ServiceHandler.cpp
 class ServiceHandler : virtual ServiceIf

TServer.cpp
 TServer(TProcessor processor,
         TServerTransport transport,
         TTransportFactory tfactory,
         TProtocolFactory pfactory)
 serve()
```

From the Thrift definition file, we generate the virtual service interface. A client class is generated, which implements the interface and uses two `TProtocol` instances to perform the I/O operations. The generated processor implements the `TProcessor` interface. The generated code has all the logic to handle RPC invocations via the `process()` call, and takes as a parameter an instance of the service interface, as implemented by the application developer.

The user provides an implementation of the application interface in separate, non-generated source code.

6.3 TServer

Finally, the Thrift core libraries provide a `TServer` abstraction. The `TServer` object generally works as follows.

- Use the `TServerTransport` to get a `TTransport`
- Use the `TTransportFactory` to optionally convert the primitive transport into a suitable application transport (typically the `TBufferedTransportFactory` is used here)
- Use the `TProtocolFactory` to create an input and output protocol for the `TTransport`
- Invoke the `process()` method of the `TProcessor` object

The layers are appropriately separated such that the server code needs to know nothing about any of the transports, encodings, or applications in play. The server encapsulates the logic around connection handling, threading, etc. while the processor deals with RPC. The only code written by the application developer lives in the definitional Thrift file and the interface implementation.

Facebook has deployed multiple `TServer` implementations, including the single-threaded `TSimpleServer`, thread-per-connection `TThreadedServer`, and thread-pooling `TThreadPoolServer`.

The `TProcessor` interface is very general by design. There is no requirement that a `TServer` take a generated `TProcessor` object. Thrift allows the application developer to easily write any type of server that operates on `TProtocol` objects (for instance, a server could simply stream a certain type of object without any actual RPC method invocation).

7. Implementation Details

7.1 Target Languages

Thrift currently supports five target languages: C++, Java, Python, Ruby, and PHP. At Facebook, we have deployed servers predominantly in C++, Java, and Python. Thrift services implemented in PHP have also been embedded into the Apache web server, providing transparent backend access to many of our frontend constructs using a `THttpClient` implementation of the `TTransport` interface.

Though Thrift was explicitly designed to be much more efficient and robust than typical web technologies, as we were designing our XML-based REST web services API we noticed that Thrift could be easily used to define our service interface. Though we do not currently employ SOAP envelopes (in the authors' opinions there is already far too much repetitive enterprise Java software to do that sort of thing), we were able to quickly extend Thrift to generate XML Schema Definition files for our service, as well as a framework for versioning different implementations of our web service. Though public web services are admittedly tangential to Thrift's core use case and design, Thrift facilitated rapid iteration and affords us the ability to quickly migrate our entire XML-based web service onto a higher performance system should the need arise.

7.2 Generated Structs

We made a conscious decision to make our generated structs as transparent as possible. All fields are publicly accessible; there are no `set()` and `get()` methods. Similarly, use of the `isset` object is not enforced. We do not include any `FieldNotSetException` construct. Developers have the option to use these fields to write more robust code, but the system is robust to the developer ignoring the `isset` construct entirely and will provide suitable default behavior in all cases.

This choice was motivated by the desire to ease application development. Our stated goal is not to make developers learn a rich new library in their language of choice, but rather to generate code that allow them to work with the constructs that are most familiar in each language.

We also made the `read()` and `write()` methods of the generated objects public so that the objects can be used outside of the context of RPC clients and servers. Thrift is a useful tool simply for generating objects that are easily serializable across programming languages.

7.3 RPC Method Identification

Method calls in RPC are implemented by sending the method name as a string. One issue with this approach is that longer method names require more bandwidth. We experimented with using fixed-size hashes to identify methods, but in the end concluded that the savings were not worth the headaches incurred. Reliably dealing with conflicts across versions of an interface definition file is impossible without a meta-storage system (i.e. to generate non-conflicting hashes for the current version of a file, we would have to know about all conflicts that ever existed in any previous version of the file).

We wanted to avoid too many unnecessary string comparisons upon method invocation. To deal with this, we generate maps from strings to function pointers, so that invocation is effectively accomplished via a constant-time hash lookup in the common case. This requires the use of a couple interesting code constructs. Because Java does not have function pointers, process functions are all private member classes implementing a common interface.

```
private class ping implements ProcessFunction {
  public void process(int seqid,
                      TProtocol iprot,
                      TProtocol oprot)
    throws TException
  { ...}
}

HashMap<String,ProcessFunction> processMap_ =
  new HashMap<String,ProcessFunction>();
```

In C++, we use a relatively esoteric language construct: member function pointers.

```
std::map<std::string,
  void (ExampleServiceProcessor::*)(int32_t,
  facebook::thrift::protocol::TProtocol*,
  facebook::thrift::protocol::TProtocol*)>
 processMap_;
```

Using these techniques, the cost of string processing is minimized, and we reap the benefit of being able to easily debug corrupt or misunderstood data by inspecting it for known string method names.

7.4 Servers and Multithreading

Thrift services require basic multithreading to handle simultaneous requests from multiple clients. For the Python and Java implementations of Thrift server logic, the standard threading libraries distributed with the languages provide adequate support. For the C++ implementation, no standard multithread runtime library exists. Specifically, robust, lightweight, and portable thread manager and timer class implementations do not exist. We investigated existing implementations, namely `boost::thread`, `boost::threadpool`, `ACE_Thread_Manager` and `ACE_Timer`.

While `boost::threads`[1] provides clean, lightweight and robust implementations of multi-thread primitives (mutexes, conditions, threads) it does not provide a thread manager or timer implementation.

`boost::threadpool`[2] also looked promising but was not far enough along for our purposes. We wanted to limit the dependency on third-party libraries as much as possible. Because `boost::threadpool` is not a pure template library and requires runtime libraries and because it is not yet part of the official Boost distribution we felt it was not ready for use in Thrift. As `boost::threadpool` evolves and especially if it is added to the Boost distribution we may reconsider our decision to not use it.

ACE has both a thread manager and timer class in addition to multi-thread primitives. The biggest problem with ACE is that it is ACE. Unlike Boost, ACE API quality is poor. Everything in ACE has large numbers of dependencies on everything else in ACE - thus forcing developers to throw out standard classes, such as STL collections, in favor of ACE's homebrewed implementations. In addition, unlike Boost, ACE implementations demonstrate little understanding of the power and pitfalls of C++ programming and take no advantage of modern templating techniques to ensure compile time safety and reasonable compiler error messages. For all these reasons, ACE was rejected. Instead, we chose to implement our own library, described in the following sections.

7.5 Thread Primitives

The Thrift thread libraries are implemented in the namespace `facebook::thrift::concurrency` and have three components:

- primitives
- thread pool manager
- timer manager

As mentioned above, we were hesitant to introduce any additional dependencies on Thrift. We decided to use `boost::shared_ptr` because it is so useful for multithreaded application, it requires no link-time or runtime libraries (i.e. it is a pure template library) and it is due to become part of the C++0x standard.

We implement standard `Mutex` and `Condition` classes, and a `Monitor` class. The latter is simply a combination of a mutex and

condition variable and is analogous to the `Monitor` implementation provided for the Java `Object` class. This is also sometimes referred to as a barrier. We provide a `Synchronized` guard class to allow Java-like synchronized blocks. This is just a bit of syntactic sugar, but, like its Java counterpart, clearly delimits critical sections of code. Unlike its Java counterpart, we still have the ability to programmatically lock, unlock, block, and signal monitors.

```
void run() {
 {Synchronized s(manager->monitor);
  if (manager->state == TimerManager::STARTING) {
    manager->state = TimerManager::STARTED;
    manager->monitor.notifyAll();
  }
 }
}
```

We again borrowed from Java the distinction between a thread and a runnable class. A `Thread` is the actual schedulable object. The `Runnable` is the logic to execute within the thread. The `Thread` implementation deals with all the platform-specific thread creation and destruction issues, while the `Runnable` implementation deals with the application-specific per-thread logic. The benefit of this approach is that developers can easily subclass the Runnable class without pulling in platform-specific super-clases.

7.6   Thread, Runnable, and shared_ptr

We use `boost::shared_ptr` throughout the `ThreadManager` and `TimerManager` implementations to guarantee cleanup of dead objects that can be accessed by multiple threads. For `Thread` class implementations, `boost::shared_ptr` usage requires particular attention to make sure `Thread` objects are neither leaked nor dereferenced prematurely while creating and shutting down threads.

Thread creation requires calling into a C library. (In our case the POSIX thread library, `libpthread`, but the same would be true for WIN32 threads). Typically, the OS makes few, if any, guarantees about when `ThreadMain`, a C thread's entry-point function, will be called. Therefore, it is possible that our thread create call, `ThreadFactory::newThread()` could return to the caller well before that time. To ensure that the returned `Thread` object is not prematurely cleaned up if the caller gives up its reference prior to the `ThreadMain` call, the `Thread` object makes a weak referenence to itself in its `start` method.

With the weak reference in hand the `ThreadMain` function can attempt to get a strong reference before entering the `Runnable::run` method of the `Runnable` object bound to the `Thread`. If no strong references to the thread are obtained between exiting `Thread::start` and entering `ThreadMain`, the weak reference returns `null` and the function exits immediately.

The need for the `Thread` to make a weak reference to itself has a significant impact on the API. Since references are managed through the `boost::shared_ptr` templates, the `Thread` object must have a reference to itself wrapped by the same `boost::shared_ptr` envelope that is returned to the caller. This necessitated the use of the factory pattern. `ThreadFactory` creates the raw `Thread` object and a `boost::shared_ptr` wrapper, and calls a private helper method of the class implementing the `Thread` interface (in this case, `PosixThread::weakRef`) to allow it to make add weak reference to itself through the `boost::shared_ptr` envelope.

`Thread` and `Runnable` objects reference each other. A `Runnable` object may need to know about the thread in which it is executing, and a Thread, obviously, needs to know what `Runnable` object it is hosting. This interdependency is further complicated because the lifecycle of each object is independent of the other. An application may create a set of `Runnable` object to be reused in different threads, or it may create and forget a `Runnable` object once a thread has been created and started for it.

The `Thread` class takes a `boost::shared_ptr` reference to the hosted `Runnable` object in its constructor, while the `Runnable` class has an explicit `thread` method to allow explicit binding of the hosted thread. `ThreadFactory::newThread` binds the objects to each other.

7.7   ThreadManager

`ThreadManager` creates a pool of worker threads and allows applications to schedule tasks for execution as free worker threads become available. The `ThreadManager` does not implement dynamic thread pool resizing, but provides primitives so that applications can add and remove threads based on load. This approach was chosen because implementing load metrics and thread pool size is very application specific. For example some applications may want to adjust pool size based on running-average of work arrival rates that are measured via polled samples. Others may simply wish to react immediately to work-queue depth high and low water marks. Rather than trying to create a complex API abstract enough to capture these different approaches, we simply leave it up to the particular application and provide the primitives to enact the desired policy and sample current status.

7.8   TimerManager

`TimerManager` allows applications to schedule `Runnable` objects for execution at some point in the future. Its specific task is to allows applications to sample `ThreadManager` load at regular intervals and make changes to the thread pool size based on application policy. Of course, it can be used to generate any number of timer or alarm events.

The default implementation of `TimerManager` uses a single thread to execute expired `Runnable` objects. Thus, if a timer operation needs to do a large amount of work and especially if it needs to do blocking I/O, that should be done in a separate thread.

7.9   Nonblocking Operation

Though the Thrift transport interfaces map more directly to a blocking I/O model, we have implemented a high performance `TNonBlockingServer` in C++ based on `libevent` and the `TFramedTransport` We implemented this by moving all I/O into one tight event loop using a state machine. Essentially, the event loop reads framed requests into `TMemoryBuffer` objects. Once entire requests are ready, they are dispatched to the `TProcessor` object which can read directly from the data in memory.

7.10   Compiler

The Thrift compiler is implemented in C++ using standard `lex/yacc` lexing and parsing. Though it could have been implemented with fewer lines of code in another language (i.e. Python Lex-Yacc (PLY) or `ocamlyacc`), using C++ forces explicit definition of the language constructs. Strongly typing the parse tree elements (debatably) makes the code more approachable for new developers.

Code generation is done using two passes. The first pass looks only for include files and type definitions. Type definitions are not checked during this phase, since they may depend upon include files. All included files are sequentially scanned in a first pass. Once the include tree has been resolved, a second pass over all files is taken that inserts type definitions into the parse tree and raises an error on any undefined types. The program is then generated against the parse tree.

Due to inherent complexities and potential for circular dependencies, we explicitly disallow forward declaration. Two Thrift structs cannot each contain an instance of the other. (Since we do not allow `null` struct instances in the generated C++ code, this would actually be impossible.)

7.11 TFileTransport

The `TFileTransport` logs Thrift requests/structs by framing incoming data with its length and writing it out to disk. Using a framed on-disk format allows for better error checking and helps with the processing of a finite number of discrete events. The `TFileWriterTransport` uses a system of swapping in-memory buffers to ensure good performance while logging large amounts of data. A Thrift log file is split up into chunks of a specified size; logged messages are not allowed to cross chunk boundaries. A message that would cross a chunk boundary will cause padding to be added until the end of the chunk and the first byte of the message are aligned to the beginning of the next chunk. Partitioning the file into chunks makes it possible to read and interpret data from a particular point in the file.

8. Facebook Thrift Services

Thrift has been employed in a large number of applications at Facebook, including search, logging, mobile, ads and the developer platform. Two specific usages are discussed below.

8.1 Search

Thrift is used as the underlying protocol and transport layer for the Facebook Search service. The multi-language code generation is well suited for search because it allows for application development in an efficient server side language (C++) and allows the Facebook PHP-based web application to make calls to the search service using Thrift PHP libraries. There is also a large variety of search stats, deployment and testing functionality that is built on top of generated Python code. Additionally, the Thrift log file format is used as a redo log for providing real-time search index updates. Thrift has allowed the search team to leverage each language for its strengths and to develop code at a rapid pace.

8.2 Logging

The Thrift `TFileTransport` functionality is used for structured logging. Each service function definition along with its parameters can be considered to be a structured log entry identified by the function name. This log can then be used for a variety of purposes, including inline and offline processing, stats aggregation and as a redo log.

9. Conclusions

Thrift has enabled Facebook to build scalable backend services efficiently by enabling engineers to divide and conquer. Application developers can focus on application code without worrying about the sockets layer. We avoid duplicated work by writing buffering and I/O logic in one place, rather than interspersing it in each application.

Thrift has been employed in a wide variety of applications at Facebook, including search, logging, mobile, ads, and the developer platform. We have found that the marginal performance cost incurred by an extra layer of software abstraction is far eclipsed by the gains in developer efficiency and systems reliability.

A. Similar Systems

The following are software systems similar to Thrift. Each is (very!) briefly described:

- *SOAP*. XML-based. Designed for web services via HTTP, excessive XML parsing overhead.
- *CORBA*. Relatively comprehensive, debatably overdesigned and heavyweight. Comparably cumbersome software installation.
- *COM*. Embraced mainly in Windows client softare. Not an entirely open solution.
- *Pillar*. Lightweight and high-performance, but missing versioning and abstraction.
- *Protocol Buffers*. Closed-source, owned by Google. Described in Sawzall paper.

Acknowledgments

References

[1] Kempf, William, "Boost.Threads", `http://www.boost.org/doc/html/threads.html`

[2] Henkel, Philipp, "threadpool", `http://threadpool.sourceforge.net`